

Assay

SPEAKER NOTES · 11 JULY 2026

The slides carry the detail. These are the words you say **over** them — the connective tissue, not the text. Don't read the screen. One idea per slide, in your own voice. Total: ~3 minutes, then demo.

SLIDE 1

On screen: "Prediction markets weight belief by money."

"Everyone here has used a prediction market. So here's the thing nobody says out loud about them: the price is set by whoever has the most money, not whoever is the most right."

Beat. Let it sit. Then move on — don't over-explain, they get it.

SLIDE 2

On screen: Assay · calibration weighting · blind replay callout

"So we weight by track record instead. Problem is, building a track record normally takes months — which is why nobody does this.

We solved that with a trick. We pull a real price chart out of our database, rip the labels off, and ask you what happens next. You can't Google it. And we already know the answer — so we score you the second you hit submit. Months becomes minutes."

If you have a chart handy, hold it up here. The blank chart sells this faster than any sentence.

SLIDE 3 — THE ONE THAT MATTERS

► SLOW DOWN. This is your whole talk.

On screen: 0.2708 / 0.2708 / 0.2505 · "not one beat a coin flip"

"We ran it on nine real people this afternoon. Here's what came back.

Weighting the crowd by calibration cut its error by seven and a half percent. On questions nobody knew the answer to. So — it works."

Pause. Then:

"But look at the middle number. That's our first weighting scheme, and it's identical to the naive one. Because every weight came out zero. Because **not one person out of nine beat a coin flip.**

Our top four all scored exactly 0.25 — that's the score you get for saying 'fifty percent' and meaning it. Everybody who was confident did worse than chance.

We didn't build an overconfidence detector on purpose. But that's what we found."

Then land it: **"The reputation layer worked. It just told us something uncomfortable about the crowd."**

"I want to be clear about what this is. It's not a market. We didn't build an AMM or an order book — the space has plenty of those.

It's one endpoint. A score any prediction market can call to find out how much to trust a given person. That's the product.

And nothing about it is financial. Same problem exists in public health, in climate, in logistics — lots of forecasters, wildly different quality, nobody keeping score. We used markets because the answer comes back in hours instead of years."

"Last thing — what's wrong with it.

Nine people is a small sample. One of our two schemes is broken. And the fix we shipped won on the live questions and lost on the holdout. That's a mixed result and I'm not going to pretend otherwise.

But it's a real system, it ran on real people, and it produced a real number in one afternoon."

Then: "**Let me show you it running.**" → go to the demo.

TAB 1 · "This is what they filled out. That's the chart. No ticker, no dates. Just — is it higher tomorrow?"

TAB 2 · "Nine forecasters. Every weight, zero. That's the coin-flip problem, right there in the table."

TAB 3 · "And that's the result. Naive crowd. Weighted crowd. Seven and a half percent better."

Stop on Tab 3. Don't click away. Let the green number be the last thing they see.

If a judge asks

ANSWERS, NOT SCRIPTS — SAY THEM HOW YOU'D SAY THEM

"Why are the charts blank?"

So you can't Google it. With a ticker and a date it's a search problem. Without them it's a judgment problem — which is the thing we're measuring.

"Price charts are random. This is unpredictable."

Completely agree, and that's the test. If it's noise, the right answer is 50% every time. But people don't say 50% — they see a downtrend and say 20%. We're not measuring whether you can predict markets. We're measuring whether you know that you can't. Our data says you don't.

"N=9 is nothing."

Correct. Which is why I'm reporting the aggregate and making no claims about any individual. The contribution is the mechanism and the cold-start solution — not the effect size.

"Your weighting didn't work on the holdout."

It didn't. It won on live and lost on holdout. At ten questions per person the signal is thin — and I'd rather show you that than only show you the half that flatters me.

"How is this different from Polymarket?"

Polymarket's product is the bet. Ours is the verdict on your judgment. Their output is a payout. Ours is a portable score that tells anyone how much to trust you. We're not competing with the market — we're the layer it should be calling.

"What's a Brier score?"

Squared error on a probability. Say 90% and be right, you score 0.01. Say 90% and be wrong, 0.81. The squaring means confident-and-wrong hurts far more than unsure-and-wrong. 0.25 is what you'd get guessing 50% on everything — that's the line nobody in our crowd crossed.

"Why forecasts instead of bets?"

A bet mixes two things — how sure you are, and how much money you have. A percentage is a clean signal of belief. That's the thing worth scoring.

"What's next?"

More questions per person — ten is too few. A non-financial question set, to test whether calibration transfers across domains. And an LLM on the leaderboard, scored identically, with a published Brier score. Every AI finance product on earth emits confident opinions with zero accountability. Ours would ship with a track record.