

Assay

SPEAKER NOTES · 11 JULY 2026

The slides carry the detail. These are the words you say **over** them — the connective tissue, not the text. Don't read the screen. One idea per slide, in your own voice. Total: ~3 minutes, then demo.

SLIDE 1

On screen: "Nobody keeps score of judgment." · late · contaminated

"Here's the objection I want to get out of the way first, because it's the one I had.

You might say: we don't need this. Reality already keeps score. The trader who's wrong blows up. The project that was never shipping in March doesn't ship in March. Time sorts it out.

True. But reality's scorecard has two problems.

It's late. The trader finds out he was overconfident after he's lost the money. That's an expensive way to learn something.

And it's dirty. A winning year mixes skill, luck, how big you sized, and whether the whole market went up. Three of those aren't judgment. P&L can't tell you which one you had.

And prediction markets don't fix this — they weight belief by money. Whoever bets most moves the price most. Rich isn't the same as right."

Beat. This is your whole setup. Don't rush it — if they buy this, they buy everything after.

SLIDE 2

On screen: Assay · "when you say 70%, are you right 70%?" · the blank chart

"So we built a clean scorecard. One question: when you say you're seventy percent sure, are you right seventy percent of the time?

No money involved, so no position sizing to muddy it. No luck, because we ask you enough times.

The catch is that a scorecard like that normally takes months to fill in — you have to wait for real events to actually happen. Which is why nobody builds this.

So here's the trick. We pull a real price chart out of the database, rip the labels off — no ticker, no dates — and ask what happens next. You can't Google it. And we already know the answer.

So we can score you the second you hit submit. Months becomes minutes."

Point at the chart on the slide. It sells this faster than any sentence you could say.

SLIDE 3 — THE ONE THAT MATTERS

► SLOW DOWN. This is your whole talk.

On screen: 0.2708 / 0.2708 / 0.2505 · "not one beat a coin flip"

"We ran it on nine real people this afternoon. Here's what came back.

Weighting the crowd by calibration cut its error by seven and a half percent. On questions nobody knew the answer to. So — it works."

Pause. Then:

"But look at the middle number. That's our first weighting scheme, and it's identical to the naive one. Because every weight came out zero. Because **not one person out of nine beat a coin flip.**

Our top four all scored exactly 0.25 — that's the score you get for saying 'fifty percent' and meaning it. Everybody who was confident did worse than chance.

We didn't build an overconfidence detector on purpose. But that's what we found."

Then land it: **"The reputation layer worked. It just told us something uncomfortable about the crowd."**

On screen: "Not a model. Not a market. A measurement layer." · the endpoint

"I want to be precise about what this is, because it gets mistaken for two things it isn't.

It's not a predictive model. It doesn't forecast anything. It has no opinion about where Bitcoin goes.

And it's not a market. Nobody bets. No odds, no money, no counterparty.

It *scores the people who do both*. Three different jobs: some people make predictions, markets turn those into a price — and this measures what each prediction was actually worth. Nobody was doing the third one.

It's one endpoint. Any market — or any trading desk, or any team — can call it and get back: how much should I trust this person?

And none of it is really about finance. Same problem on a trading desk, in a hospital, in a project plan. Lots of confident people, wildly different quality, nobody keeping score. We used markets because the answer comes back in hours instead of years."

On screen: Built / limitations / next

"Last thing — what's wrong with it. And I want to be blunt.

We tested the weighting twice. On the live questions, it helped — that's the seven and a half percent. Then we tested it on a set of questions the system had never seen, and it **hurt**. Made the crowd worse.

Same idea. Same nine people. Opposite results.

Which means the honest answer is: I don't have enough data to know if this works yet. Nine people, ten questions — a real effect would win both tests. A dead idea would lose both. Winning one and losing one is just noise.

So I'm showing you both numbers instead of only the one that makes me look good.

What I do know: it's a real system, it ran on real people, and it found something true about them in one afternoon."

Then: "**Let me show you it running.**" → go to the demo.

TAB 1 · "This is what they filled out. That's the chart. No ticker, no dates. Just — is it higher tomorrow?"

TAB 2 · "Nine forecasters. Every weight, zero. That's the coin-flip problem, right there in the table."

TAB 3 · "And that's the result. Naive crowd. Weighted crowd. Seven and a half percent better."

Stop on Tab 3. Don't click away. Let the green number be the last thing they see.

If a judge asks

ANSWERS, NOT SCRIPTS — SAY THEM HOW YOU'D SAY THEM

"Why are the charts blank?"

So you can't Google it. With a ticker and a date it's a search problem. Without them it's a judgment problem — which is the thing we're measuring.

"Price charts are random. This is unpredictable."

Completely agree, and that's the test. If it's noise, the right answer is 50% every time. But people don't say 50% — they see a downtrend and say 20%. We're not measuring whether you can predict markets. We're measuring whether you know that you can't. Our data says you don't.

"Is this a predictive model or a market?"

Neither. It doesn't predict anything — it has no view on where any price goes. And nobody bets: no odds, no money, no counterparty. It's a measurement layer. Three separate jobs: people make predictions, markets turn those into a price, and this scores what each prediction was worth. We built the third one.

"Why does this matter? Traders find out if they're good over time anyway."

They do — and they pay for it. The feedback from reality is accurate and late; you learn you were overconfident by losing the money. This gets you the same answer before the bill arrives. And P&L can't isolate judgment anyway — it mixes skill, luck, and how big you sized. Calibration strips all that out. The real kicker: in principle anyone could audit their own track record. In practice nobody does. Nine people here just found out they'd never been calibrated, and not one of them knew.

"N=9 is nothing."

Correct. Which is why I'm reporting the aggregate and making no claims about any individual. The contribution is the mechanism and the cold-start solution — not the effect size.

"Your weighting didn't work on the holdout."

Right — it won one test and lost the other. Which tells you I don't have enough data yet, not that the idea is wrong. If the effect were real it would win both; if it were dead it would lose both. Winning one and losing one is what noise looks like at nine people. The fix is more forecasters and more questions, and I'd rather tell you that than show you only the half that flatters me.

"How is this different from Polymarket?"

Polymarket's product is the bet. Ours is the verdict on your judgment. Their output is a payout. Ours is a portable score that tells anyone how much to trust you. We're not competing with the market — we're the layer it should be calling.

"What's a Brier score?"

Squared error on a probability. Say 90% and be right, you score 0.01. Say 90% and be wrong, 0.81. The squaring means confident-and-wrong hurts far more than unsure-and-wrong. 0.25 is what you'd get guessing 50% on everything — that's the line nobody in our crowd crossed.

"Why forecasts instead of bets?"

A bet mixes two things — how sure you are, and how much money you have. A percentage is a clean signal of belief. That's the thing worth scoring.

"What's next?"

More questions per person — ten is too few. A non-financial question set, to test whether calibration transfers across domains. And an LLM on the leaderboard, scored identically, with a published Brier score. Every AI finance product on earth emits confident opinions with zero accountability. Ours would ship with a track record.